



# Real Time Data Streaming with **Apache Spark**

Arjun Yadav (309939)  
Bhomit Kapoor (310007)  
Drashti Pandya (310112)  
Karan Sethi (309984)



## Appendix

Sr. No.	Content	Page No.
1.	Abstract	2
2.	Steps for initiating project	2
3.	API data	5
4.	Cloud deployment	6
5.	Output window for Ubuntu instance	9
6.	Conclusion	12

## **Abstract**

Data today are being manipulated from multiple resources which are huge in size and complex as it has incremental addition with existing database. Because data are just not only piece of information but with that we do Analytics too, we shall define a unified or defined repository for Data. In today's IT (Information Technology) Age the data created by any business in faster then ever before, creation of new data is tremendous. To be specific this data is known as LIVE DATA. Live Data say for instance, Flight Data, Social Media feeds, News Reports, Sports Score Data, Satellite Information, Shipping or Logistics domain, Stock Exchange, Telecom Industries Information Transformation, Weather Forecasting, Search Engine Search Repository and so on, they are all supplied within an API. Within this Project our team will work on various Tools, Software's, Technologies and Infrastructure to store the collective information in a console and deploy it over cloud as well as on premise. For developing this project on Live Data Streaming we are using Kafka, Docker, Kubernetes Cluster, Spark, Java, Scala, Zookeeper here all the mentioned technologies have its significant role. API is a key resource for this assignment here we are using Alpha-Vantage API, It will allow us to gather the information for Stock Exchange Market. We are using Kafka to rectify the information/data gathered from API as Source API will bring all information but useful information for our project perspective will be then structured with help of Kafka. Kafka will debug the error and organize the data in parameterized manner. Data from Kafka will then be outputted in JSON format. Now that we have data we will create a pathway or pipeline to Stream data which is key objective of assignment, we will be using Apache Spark to stream Realtime Data. This data will be shipped or wrapped inside a Docker Container and it will be Hosted to a cloud Platform.

## **Steps for Initiating Project.**

### **Start Zookeeper**

1. go to zookeeper file
2. C:\zookeeper-3.5.6>zkserver run....

### **Start Kafka**

1. go to kafka file
2. .\bin\windows\kafka-server-start.bat .\config\server.properties run code

### **Create Topic**

1. Open command prompt and go to Apache Kafka installation directory.
2. Go to \bin\windows directory.
3. Run the following command.
4. kafka-topics.bat --create --zookeeper localhost:2181 --replication-factor 1 -- partitions 1 --topic sql-insert

**Create producer**

```
C:\kafka_2.11-2.3.1\bin\windows>kafka-console-producer.bat --broker-list localhost:9092 --  
topic Kafka_Example
```

**Create Consumer**

```
C:\kafka_2.11-2.3.1\bin\windows>kafka-console-consumer.bat --bootstrap-server  
localhost:9092 --topic Kafka_Example --from-beginning
```

**Create config folder**

- create KafkaConfig file
- initialize all config in producerFactory
- config bootstrap, key and value

**Create model folder**

- create User class
- define variable
- set get and set method

**Create resource folder**

- create UserResource class
- KafkaTemplate with object
- @GetMapping to get data from JSON
- use Topic name which already created

**To run Spark**

- open cmd
- run **spark-shell**

## create performance.scala

### create object file

```
Performance.scala x Stock-Analysis-and-Prediction.iml x
No Scala SDK in module Setup Scala SDK

13 def
14 def performanceMeasure(sc: SparkContext, sqlContext: SQLContext, name: String): Unit = {
15   val filename = name.split( regex = "." ).toSeq
16   val business = sqlContext.read
17     .format( source = "csv" )
18     .option( "header", "true" )
19     .load( System.getProperty( "perf.dir" ) + "/" + filename(0) + "_fundamentals.csv" )
20
21   business.toDF().registerTempTable( tableName = "new" )
22
23   val spark_Data = sqlContext.sql( sqlText = "Select `Cash and Cash Equivalents` / `Total Current Liabilities` As Cash_Ratio, `Total Cu"
24
25   spark_Data.toDF().registerTempTable( tableName = "required_data" )
26
27   val bad_performance = sqlContext.sql( sqlText = "Select * from required_data WHERE Debt_Ratio >=2 AND Cash_Ratio < 1 AND Current_Ra"
28
29   val good_performance = sqlContext.sql( sqlText = "Select * from required_data WHERE Debt_Ratio >=2 AND Cash_Ratio < 1 AND Current_f"
30
31   bad_performance.coalesce( numPartitions = 1 ).write.format( source = "csv" )
32     .save( System.getProperty( "perf.dir" ) + "/" + filename(0) + "_badPerfData.csv" )
33
34   good_performance.coalesce( numPartitions = 1 ).write.format( source = "csv" )
35     .save( System.getProperty( "perf.dir" ) + "/" + filename(0) + "_goodPerfData.csv" )
36 }
```

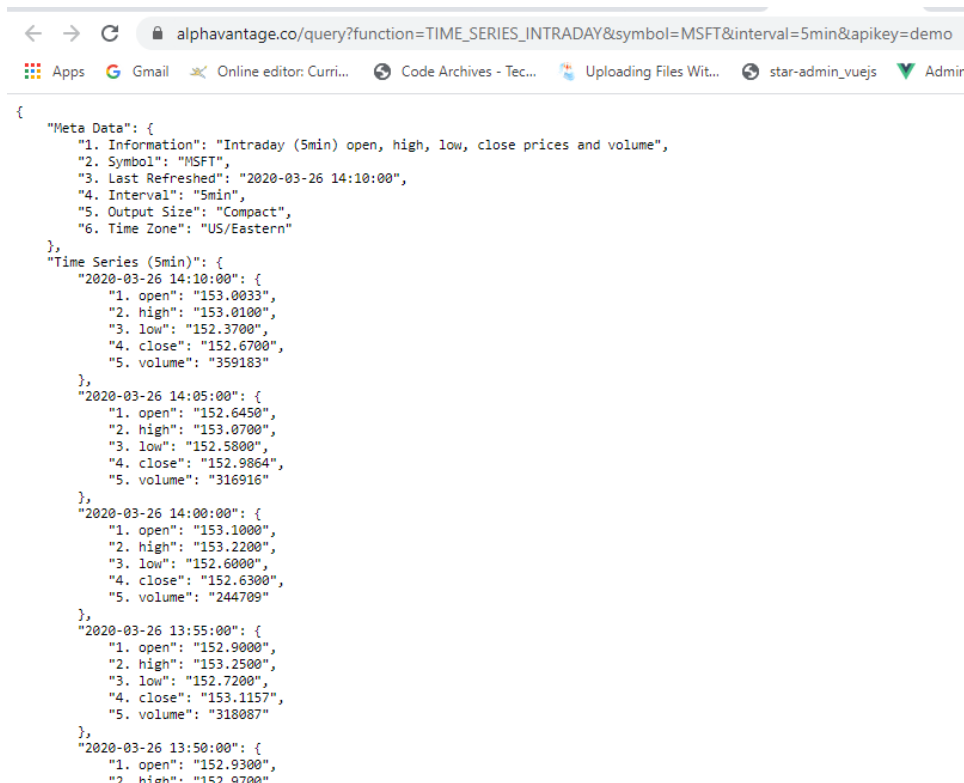
## Create prediction.scala

```
Performance.scala x Prediction.scala x Stock-Analysis-and-Prediction.iml x
No Scala SDK in module Setup Scala SDK

66 def getYahooFinData(sc: SparkContext, sqlContext: SQLContext, name: String): Unit = {
67   import sqlContext.implicits._
68   val filename = name.split( regex = "." ).toSeq
69   val fin_data = sqlContext.read
70     .format( source = "csv" )
71     .option( "header", "true" )
72     .load( System.getProperty( "pred.dir" ) + "/" + name )
73
74   val wSpec1 = Window.orderBy( colName = "Date" )
75   val fin_data1 = fin_data.withColumn( colName = "YC", lag( fin_data( "Close" ), offset = 1 ).over( wSpec1 ) )
76   val fin_data2 = fin_data1.withColumn( colName = "TDAC", lag( fin_data1( "Close" ), offset = 10 ).over( wSpec1 ) )
77
78   val wSpec2 = Window.orderBy( colName = "Date" ).rowsBetween( -4, 0 )
79   val wSpec3 = Window.orderBy( colName = "Date" ).rowsBetween( -9, 0 )
80   val fin_data3 = fin_data2.withColumn( colName = "Lowest Low", min( fin_data2( "Low" ) ).over( wSpec2 ) )
81   val fin_data4 = fin_data3.withColumn( colName = "Highest High", max( fin_data3( "High" ) ).over( wSpec2 ) )
82   val fin_data5 = fin_data4.withColumn( colName = "MA", avg( fin_data4( "Close" ) ).over( wSpec3 ) )
83
84   val filteredRdd1 = fin_data5.rdd.zipWithIndex().collect { case (r, i) if i > 9 => r }
85
86   val newDf1 = sqlContext.createDataFrame( filteredRdd1, fin_data5.schema )
87   val filteredRdd = newDf1.rdd.zipWithIndex().collect { case (r, i) if i < 400 => r }
88   val newDf = sqlContext.createDataFrame( filteredRdd, fin_data5.schema )
89
90   val fin_data6 = newDf.withColumn( colName = "Volume1", newDf( "Volume" ).cast( org.apache.spark.sql.types.LongType ) )
}
```

## API data

### in JSON format



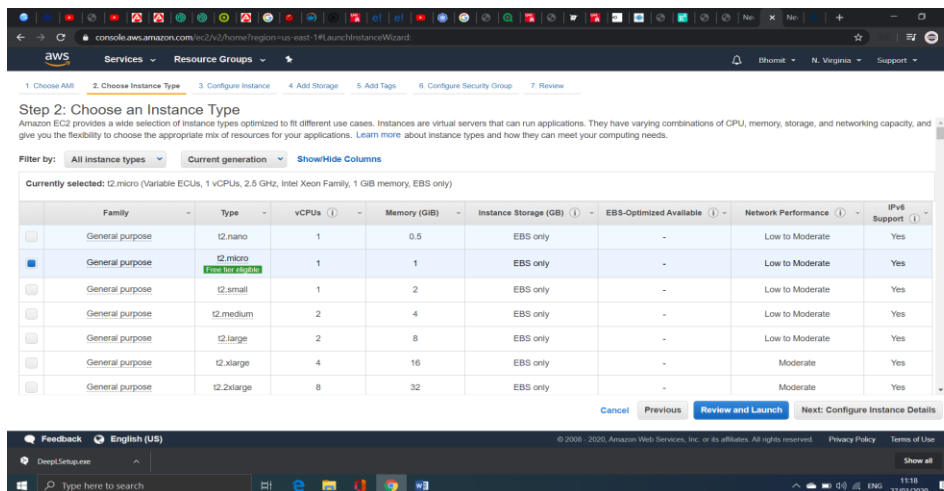
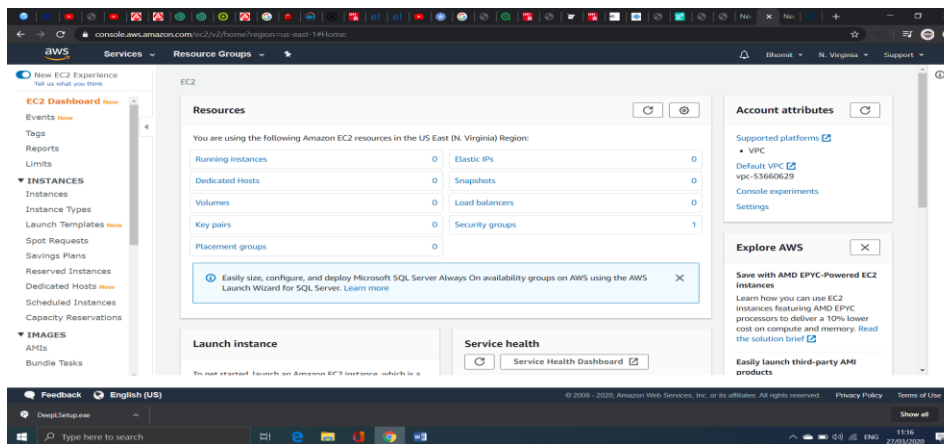
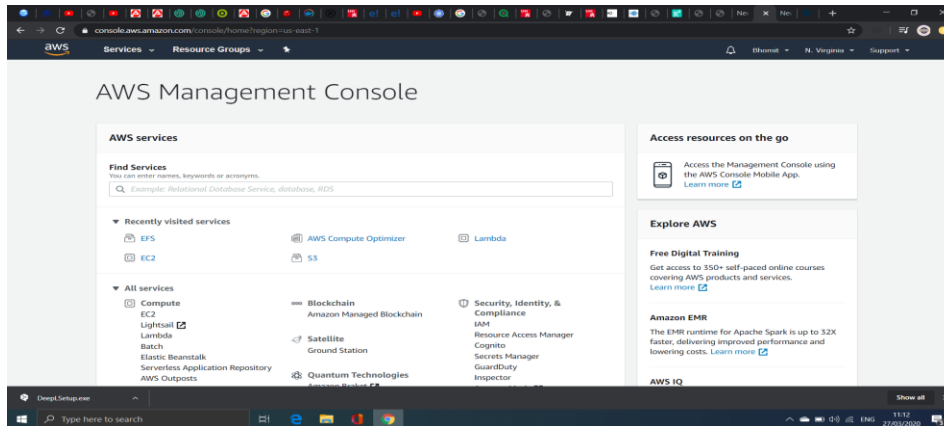
The screenshot shows a web browser with the URL `alphavantage.co/query?function=TIME_SERIES_INTRADAY&symbol=MSFT&interval=5min&apikey=demo`. The browser's taskbar at the bottom shows several open applications: Apps, Gmail, Online editor: Curri..., Code Archives - Tec..., Uploading Files Wit..., star-admin\_vuejs, and Admin. The main content area displays a JSON response from the API. The JSON is structured as follows:

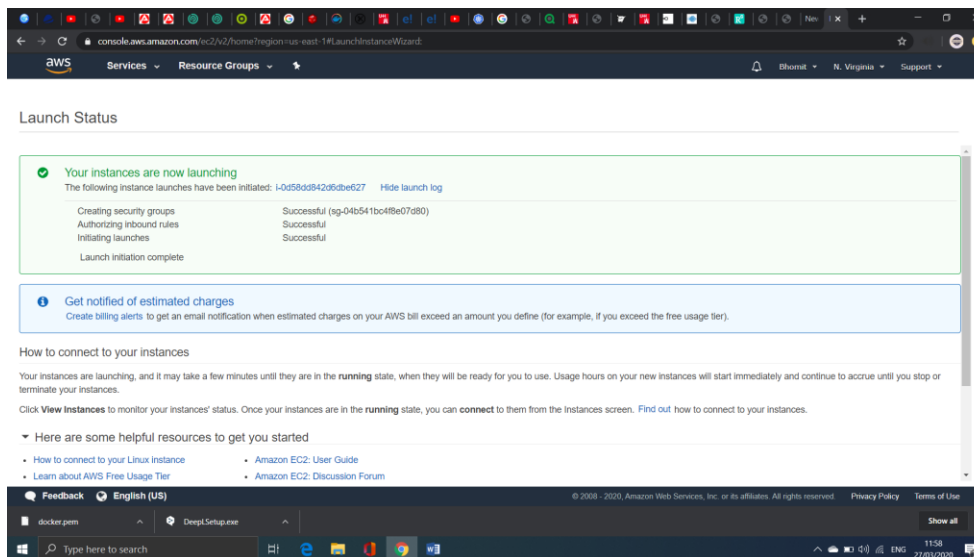
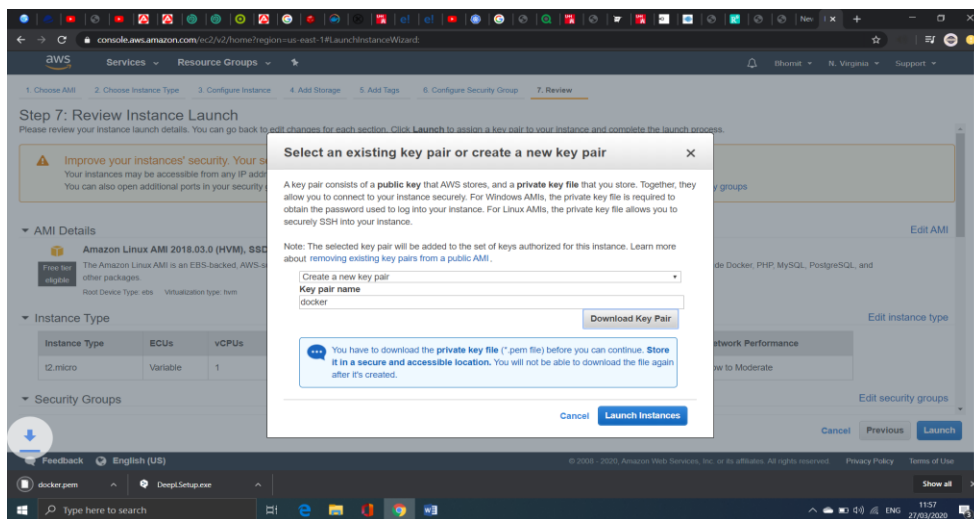
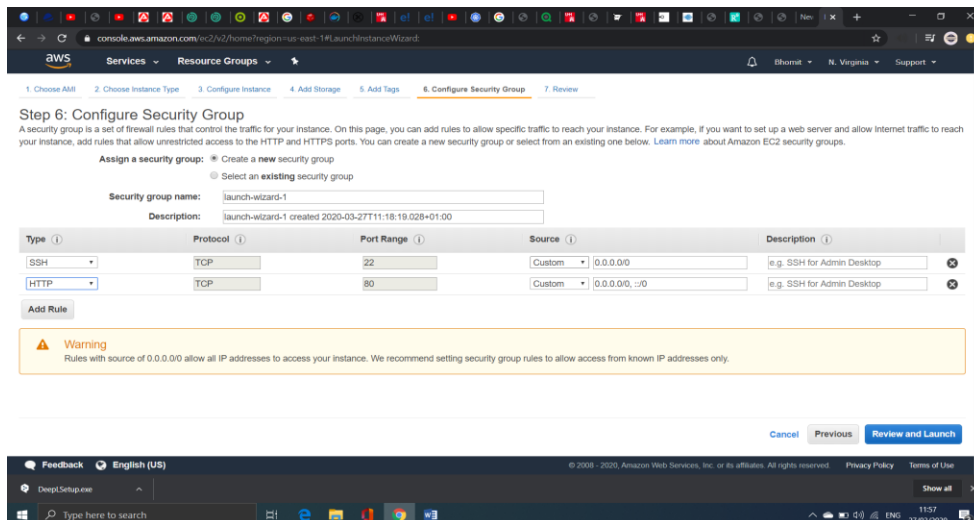
```
{
  "Meta Data": {
    "1. Information": "Intraday (5min) open, high, low, close prices and volume",
    "2. Symbol": "MSFT",
    "3. Last Refreshed": "2020-03-26 14:10:00",
    "4. Interval": "5min",
    "5. Output Size": "Compact",
    "6. Time Zone": "US/Eastern"
  },
  "Time Series (5min)": {
    "2020-03-26 14:10:00": {
      "1. open": "153.0033",
      "2. high": "153.0100",
      "3. low": "152.3700",
      "4. close": "152.6700",
      "5. volume": "359183"
    },
    "2020-03-26 14:05:00": {
      "1. open": "152.6450",
      "2. high": "153.0700",
      "3. low": "152.5800",
      "4. close": "152.9864",
      "5. volume": "316916"
    },
    "2020-03-26 14:00:00": {
      "1. open": "153.1000",
      "2. high": "153.2200",
      "3. low": "152.6000",
      "4. close": "152.6300",
      "5. volume": "244709"
    },
    "2020-03-26 13:55:00": {
      "1. open": "152.9000",
      "2. high": "153.2500",
      "3. low": "152.7200",
      "4. close": "153.1157",
      "5. volume": "318087"
    },
    "2020-03-26 13:50:00": {
      "1. open": "152.9300",
      "2. high": "152.9700",
      "3. low": "152.9300",
      "4. close": "152.9300",
      "5. volume": "10000"
    }
  }
}
```

With Alpha-Vantage API we can extract weekly data, monthly data, yearly data, intermediate data etc. apart from this we use some function for 5min, and 1 min duration changes in stock data. Above is an example of 5 min duration of data. This API returns inter day time series (timestamp, high, low, open, close, volume) of the equity specific data. We can get JSON and CSV file to collect and retrieve data from Alphavantage. The most beneficial thing is it gives prices and volume information of the current trading day, update Realtime. We can also use 20+ year historic data to stream and compare with real time.


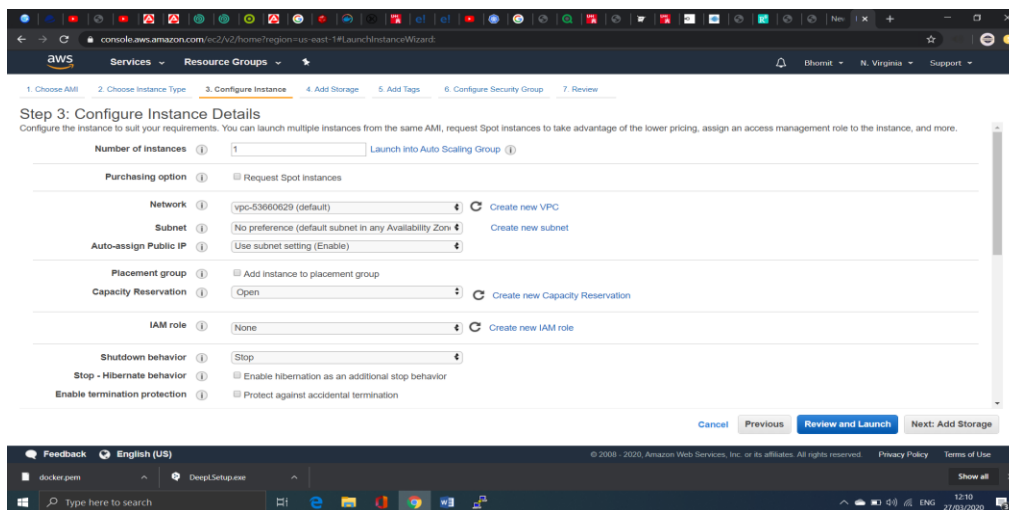
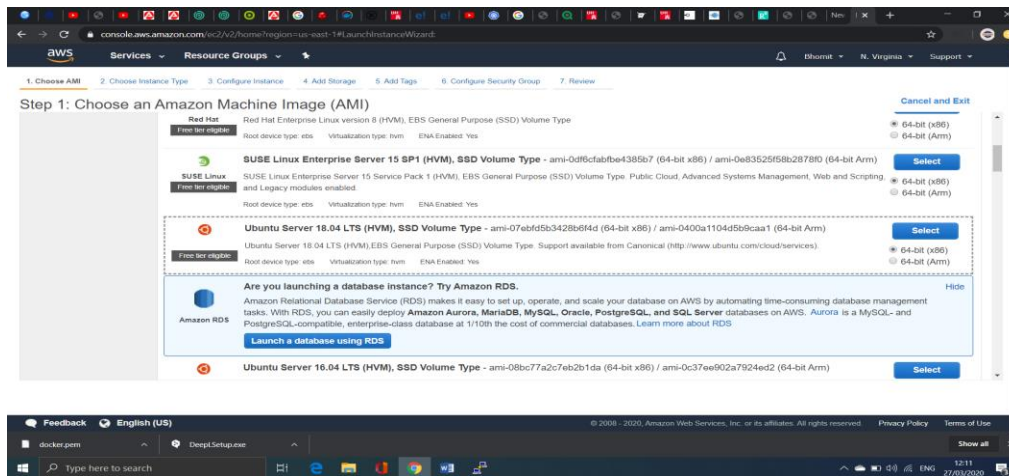
# Cloud Deployment

## AWS Console Cloud









## Output window for Ubuntu instance

```
ubuntu@ip-172-31-85-201:~$
login as: ubuntu
Authenticating with public key "imported-openssh-key"
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1057-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Fri Mar 27 11:21:39 UTC 2020

System load: 0.66      Processes:           89
Usage of /: 13.6% of 7.69GB Users logged in:       0
Memory usage: 15%      IP address for eth0: 172.31.85.201
Swap usage: 0%

0 packages can be updated.
0 updates are security updates.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$
```

```
ubuntu@ip-172-31-85-201:~$
login as: ubuntu
Authenticating with public key "imported-openssh-key"
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1057-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Fri Mar 27 11:21:39 UTC 2020

System load: 0.66      Processes:           89
Usage of /: 13.6% of 7.69GB Users logged in:       0
Memory usage: 15%      IP address for eth0: 172.31.85.201
Swap usage: 0%

0 packages can be updated.
0 updates are security updates.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$ sudo snap install docker
Fetch and check assertions for snap "docker" (423)
```

```
ubuntu@ip-172-31-85-201:~$
login as: ubuntu
Authenticating with public key "imported-openssh-key"
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-1057-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

System information as of Fri Mar 27 11:21:39 UTC 2020

System load: 0.66      Processes:           89
Usage of /: 13.6% of 7.69GB Users logged in:       0
Memory usage: 15%      IP address for eth0: 172.31.85.201
Swap usage: 0%

0 packages can be updated.
0 updates are security updates.

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$ sudo snap install docker
docker 18.09.9 from Canonical/ installed
ubuntu@ip-172-31-85-201:~$
```

```
ubuntu@ip-172-31-85-201:~$
The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$ sudo snap install docker
docker 18.09.9 from Canonical/ installed
ubuntu@ip-172-31-85-201:~$ sudo docker version
Client:
 Version:      18.09.9
 API version:  1.39
 Go version:   go1.13.4
 Git commit:   1752eb3
 Built:        Sat Nov 16 01:05:26 2019
 OS/Arch:      linux/amd64
 Experimental:  false

Server:
 Engine:
  Version:      18.09.9
  API version:  1.39 (minimum version 1.12)
  Go version:   go1.13.4
  Git commit:   9552f2b
  Built:        Sat Nov 16 01:07:48 2019
  OS/Arch:      linux/amd64
  Experimental:  false
ubuntu@ip-172-31-85-201:~$ sudo docker run -d tomcat
Unable to find image 'tomcat:latest' locally
latest: Pulling from library/tomcat
30e431f79093: Extracting [====>] 2.621MB/50.38MB
d88c6d374ea5: Download complete
c85513200d84: Download complete
35769680e827: Downloading [=====] 42.82MB/51.79MB
e27ce2095ec2: Download complete
3943eeae6b7c: Download complete
3ed8c8ae72a6: Downloading [=====] 47.49MB/104.2MB
91d1e510d72b: Download complete
415cc4506e71: Download complete
a79d88064227: Download complete
```

```
ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$ sudo snap install docker
docker 18.09.9 from Canonical✓ installed
ubuntu@ip-172-31-85-201:~$ sudo docker version
Client:
 Version:      18.09.9
 API version:  1.39
 Go version:   go1.13.4
 Git commit:   1752eb3
 Built:        Sat Nov 16 01:05:26 2019
 OS/Arch:      linux/amd64
 Experimental: false

Server:
 Engine:
  Version:      18.09.9
  API version:  1.39 (minimum version 1.12)
  Go version:   go1.13.4
  Git commit:   9552f2b
  Built:        Sat Nov 16 01:07:48 2019
  OS/Arch:      linux/amd64
  Experimental: false
ubuntu@ip-172-31-85-201:~$ sudo docker run -d tomcat
Unable to find image 'tomcat:latest' locally
latest: Pulling from library/tomcat
50e431f79093: Pull complete
dd8c6d374ea5: Pull complete
c85513200d84: Pull complete
c85513200d84: Pull complete
55769680e827: Pull complete
e27ce2095ec2: Pull complete
5943eeae6b7c: Pull complete
3ed8ceae72a6: Pull complete
91d1e510d72b: Pull complete
415cc4506e71: Pull complete
a79d8064227: Pull complete
Digest: sha256:b707d3b8b4f40951ca2f387c24ab9f78800c69c90740f0cca937a1b95204b3a4
Status: Downloaded newer image for tomcat:latest
e8a92ec60b582a2a0264618612f1adfe5f98570f8996622ab919853cfd26211
ubuntu@ip-172-31-85-201:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS               NAMES
e8a92ec60b58       tomcat              "catalina.sh run"       25 seconds ago     Up 23 seconds      8080/tcp            hopeful_driscoll
ubuntu@ip-172-31-85-201:~$
```

```
ubuntu@ip-172-31-85-201:~$ sudo docker version
sudo: docker: command not found
ubuntu@ip-172-31-85-201:~$ sudo snap install docker
docker 18.09.9 from Canonical✓ installed
ubuntu@ip-172-31-85-201:~$ sudo docker version
Client:
 Version:      18.09.9
 API version:  1.39
 Go version:   go1.13.4
 Git commit:   1752eb3
 Built:        Sat Nov 16 01:05:26 2019
 OS/Arch:      linux/amd64
 Experimental: false

Server:
 Engine:
  Version:      18.09.9
  API version:  1.39 (minimum version 1.12)
  Go version:   go1.13.4
  Git commit:   9552f2b
  Built:        Sat Nov 16 01:07:48 2019
  OS/Arch:      linux/amd64
  Experimental: false
ubuntu@ip-172-31-85-201:~$ sudo docker run -d tomcat
Unable to find image 'tomcat:latest' locally
latest: Pulling from library/tomcat
50e431f79093: Pull complete
dd8c6d374ea5: Pull complete
c85513200d84: Pull complete
c85513200d84: Pull complete
55769680e827: Pull complete
e27ce2095ec2: Pull complete
5943eeae6b7c: Pull complete
3ed8ceae72a6: Pull complete
91d1e510d72b: Pull complete
415cc4506e71: Pull complete
a79d8064227: Pull complete
Digest: sha256:b707d3b8b4f40951ca2f387c24ab9f78800c69c90740f0cca937a1b95204b3a4
Status: Downloaded newer image for tomcat:latest
e8a92ec60b582a2a0264618612f1adfe5f98570f8996622ab919853cfd26211
ubuntu@ip-172-31-85-201:~$ sudo docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS               NAMES
e8a92ec60b58       tomcat              "catalina.sh run"       25 seconds ago     Up 23 seconds      8080/tcp            hopeful_driscoll
ubuntu@ip-172-31-85-201:~$
```

```
ubuntu@ip-172-31-85-201:~$ cat /var/run/docker.sock | nc -l 172.31.85.201
E at unix:///var/run/docker.sock: Post http://172.31.85.201:2376/v1.39/c
ontainers/create?name=m1: dial unix /var/run/docker.sock: connect: permission de
nied.
See 'docker run --help'.
ubuntu@ip-172-31-85-201:~$ uname -r
4.15.0-1057-aws
ubuntu@ip-172-31-85-201:~$ sudo curl -sSL https://get.docker.com/ | sh
# Executing docker install script, commit: 442e66405c304fa92af8aadaa1d9b31bf4b0a
d94
Warning: the "docker" command appears to already exist on this system.

If you already have Docker installed, this script can cause trouble, which is
why we're displaying this warning and provide the opportunity to cancel the
installation.

If you installed the current Docker package using this script and are using it
again to update Docker, you can safely ignore this message.

You may press Ctrl+C now to abort this script.
^C sleep 20
^C
ubuntu@ip-172-31-85-201:~$ sudo docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b9308010825: Pull complete
Digest: sha256:f9dfdf6363d84ef479d645ab588515eae030f611a56f3a7ac7f2fdd86d7e4e
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (amd64)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
ubuntu@ip-172-31-85-201:~$
```

That's all we need to get Nginx up! Paste the IP address of your Droplet into a web browser and you should see Nginx's "Welcome to nginx!" page.

You'll also notice in your shell session that the log for Nginx is being updated when you make requests to your server, because we're running our container interactively.

Let's hit the break shortcut **CTRL+C** to get back to our shell session.

If you try to load the page now, you'll get a "connection refused" page. This is because we shut down our container. We can verify this with this command:

```
$ sudo docker ps -a
```

You should see something similar to the output shown below.

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
05012ab02ca1	nginx	"nginx -g 'daemon off'"	57 seconds ago	Exited (0)

We can see that our Docker container has exited.

Sign up for our newsletter.  
Get the latest tutorials on SysAdmin and open source topics.

Enter your email address

```
ubuntu@ip-172-31-85-201:~$ sudo docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS
e95eb51ecf82        nginx              "nginx -g 'daemon of-" 27 seconds ago      Exited (0) 5 seconds ago
a9d1f520e7d4        docker-nginx       "/hello"            3 minutes ago       Exited (0) 3 minutes ago
e6a22ec60b56        elegant_beaver     "catalina.sh run"    34 minutes ago      Up 34 minutes      8080/tcp
ubuntu@ip-172-31-85-201:~$
```

Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
The following additional packages will be installed:  
 bridge-utils cgroupfs-mount contained pipz runc ubuntu-fan  
Suggested packages:  
 ifupdown aufs-tools deboststrap docker-doc rinse rfs-fuse | zfsutils  
The following new packages will be installed:  
 bridge-utils cgroupfs-mount contained docker.io pipz runc ubuntu-fan  
0 upgraded, 7 newly installed, 0 to remove and 72 not upgraded.  
Need to get 63.6 MB of archives.  
After this operation, 319 MB of additional disk space will be used.  
Do you want to continue? [Y/n] Y  
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic/universe amd64 pipz amd64 2.4-1 [57.4 kB]  
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic/main amd64 bridge-utils amd64 1.5-15ubuntu1 [30.1 kB]  
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic/universe amd64 cgroupfs-mount all 1.4 [62.0 B]  
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic-updates/universe amd64 runc amd64 1.0.0-rc10-0ubuntu1-18.04.2 [2009 kB]  
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic-updates/universe amd64 contained amd64 1.3.3-0ubuntu1-18.04.1 [21.7 MB]  
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic-updates/universe amd64 docker.io amd64 19.03.6-0ubuntu1-18.04.1 [39.9 MB]  
Get:7 http://us-east-1.ec2.archive.ubuntu.com/ubuntu bionic/main amd64 ubuntu-fan all 0.12.10 [34.7 kB]  
Fetched 63.6 MB in 1s (69.9 MB/s)  
Preconfiguring packages ...  
Selecting previously unselected package pipz.  
(Reading database ... 56358 files and directories currently installed.)  
Preparing to unpack .../0-pipz\_2.4-1\_amd64.deb ...  
Unpacking pipz (2.4-1) ...  
Selecting previously unselected package bridge-utils.  
Preparing to unpack .../1-bridge-utils\_1.5-15ubuntu1\_amd64.deb ...  
Unpacking bridge-utils (1.5-15ubuntu1) ...  
Selecting previously unselected package cgroupfs-mount.  
Preparing to unpack .../2-cgroupfs-mount\_1.4\_all.deb ...  
Unpacking cgroupfs-mount (1.4) ...  
Selecting previously unselected package runc.  
Preparing to unpack .../3-runc\_1.0.0-rc10-0ubuntu1-18.04.2\_amd64.deb ...  
Unpacking runc (1.0.0-rc10-0ubuntu1-18.04.2) ...  
Selecting previously unselected package contained.  
Preparing to unpack .../4-contained\_1.3.3-0ubuntu1-18.04.1\_amd64.deb ...  
Unpacking contained (1.3.3-0ubuntu1-18.04.1) ...  
Selecting previously unselected package docker.  
Preparing to unpack .../5-docker.io\_19.03.6-0ubuntu1-18.04.1\_amd64.deb ...  
Unpacking docker.io (19.03.6-0ubuntu1-18.04.1) ...  
Selecting previously unselected package ubuntu-fan.  
Preparing to unpack .../6-ubuntu-fan\_0.12.10\_all.deb ...  
Unpacking ubuntu-fan (0.12.10) ...  
Setting up runc (1.0.0-rc10-0ubuntu1-18.04.2) ...  
Setting up cgroupfs-mount (1.4) ...  
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /lib/systemd/system/containerd.service.  
Setting up bridge-utils (1.5-15ubuntu1) ...  
Setting up ubuntu-fan (0.12.10) ...  
Created symlink /etc/systemd/system/multi-user.target.wants/ubuntu-fan.service → /lib/systemd/system/ubuntu-fan.service.  
Setting up pipz (2.4-1) ...  
Setting up docker.io (19.03.6-0ubuntu1-18.04.1) ...  
Adding group 'docker' (GID 115) ...  
Done.  
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /lib/systemd/system/docker.socket.  
docker.service is a disabled or a static unit, not starting it.  
Processing triggers for systemd (237-3ubuntu1.33) ...  
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...  
Processing triggers for ureadahead (0.100.0-21) ...  
ubuntu@ip-172-31-85-201:~\$ sudo systemctl start docker  
Failed to enable unit: Unit file docker.service does not exist.  
ubuntu@ip-172-31-85-201:~\$ sudo systemctl enable docker  
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /lib/systemd/system/docker.service.  
ubuntu@ip-172-31-85-201:~\$

## Enable Docker

```
ubuntu@ip-172-31-85-201:~$ kubectl
kubectl
  kubeconfig  Manage the kubeconfig file
  rollout     Manage the rollout of a resource
  scale       Set a new size for a Deployment, ReplicaSet or Replication Controller
  autoscale   Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:
  certificate Modify certificate resources.
  cluster-info Display cluster info
  top         Display Resource (CPU/Memory/Storage) usage.
  cordon      Mark node as unschedulable
  uncordon    Mark node as schedulable
  drain       Drain node in preparation for maintenance
  taint       Update the taints on one or more nodes

Troubleshooting and Debugging Commands:
  describe    Show details of a specific resource or group of resources
  logs        Print the logs for a container in a pod
  attach      Attach to a running container
  exec        Execute a command in a container
  port-forward Forward one or more local ports to a pod
  proxy       Run a proxy to the Kubernetes API server
  cp          Copy files and directories to and from containers.
  auth        Inspect authorization

Advanced Commands:
  diff        Diff live version against would-be applied version
  apply       Apply a configuration to a resource by filename or stdin
  patch       Update field(s) of a resource using strategic merge patch
  replace     Replace a resource by filename or stdin
  wait        Experimental: Wait for a specific condition on one or many resources.
  convert     Convert config files between different API versions
  customize   Build a kustomization target from a directory or a remote url.

Settings Commands:
  label       Update the labels on a resource
  annotate    Update the annotations on a resource
  completion  Output shell completion code for the specified shell (bash or zsh)

Other Commands:
  api-resources Print the supported API resources on the server
  api-versions  Print the supported API versions on the server, in the form of "group/version"
  config       Modify kubeconfig files
  plugin       Provides utilities for interacting with plugins.
  version      Print the client and server version information

Usage:
  kubectl [flags] [options]

Use "kubectl <command> --help" for more information about a given command.
Use "kubectl options" for a list of global command-line options (applies to all commands).
```

```
ubuntu@ip-172-31-85-201:~$ kubectl version
Client Version: version.Info{Major:"1", Minor:"17", GitVersion:"v1.17.3", GitCommit:"06ad960bdfcd3b39c8310aa92d1e7c12ce610213", GitTreeState:"clean", BuildDate:"2020-02-12T13:43:46Z", GoVersion:"go1.13.7", Compiler:"gc", Platform:"linux/amd64"}
The connection to the server localhost:8080 was refused - did you specify the right host or port?
ubuntu@ip-172-31-85-201:~$
```

## Conclusion

By utilizing the technologies such as Kafka, Spark, Docker, Kubernetes. We were able to achieve the key objective of Streaming Live Data. The data we got was from Alpha-Vantage API for real time as well as historical data of stock market. The output is getting stored to the docker container which than will be orchestrated by kubernetes clusters.